
DEME–ErisML Governance Plugin for Gazebo
via DEME MCP Server in a Medical Triage Emergency Scenario

1. Abstract

This whitepaper describes the **DEME–ErisML Governance Plugin for Gazebo**, a real-time ethical oversight layer for autonomous agents operating in the Gazebo robotics simulator. The plugin connects Gazebo simulations to a **DEME MCP (Moral Control Plane) server** running the **ErisML** library, which implements DEME’s multi-dimensional *moral landscape* framework.

We focus on a “**Medical Triage Emergency**” scenario in which an autonomous triage agent must allocate scarce critical-care resources among multiple patients under time pressure. The plugin intercepts candidate actions in the simulation, converts state into structured **EthicalFacts**, queries the DEME MCP server, and enforces **veto**es and ranked choices in real time. This enables researchers to:

- Prototype governance profiles (e.g., hospital ethics policies),
- Test learning agents under hard ethical constraints,
- Generate tamper-evident decision logs for analysis and audit.

The result is a complete loop from **policy definition** to **ethically governed behavior** in a realistic robotics simulation environment.

2. Background and Motivation

2.1 DEME and ErisML

DEME 2.0 treats machine ethics as a first-class engineering subsystem. Each candidate action is mapped to a vector in a **moral landscape**—dimensions such as expected physical harm, rights respect, fairness, autonomy, legitimacy and epistemic quality. A **governance profile** then interprets this landscape through:

- Hard **veto regions** (non-negotiable constraints),
- **Lexicographic priorities** (e.g., minimize harm before maximizing fairness),
- A **scalarization function** for ranking permissible options.

ErisML is a software library that operationalizes this framework:

- Transforms **EthicalFacts** into moral vectors,
- Applies governance profiles,
- Returns veto decisions, scores and explanatory reason codes.

2.2 Gazebo and the need for ethical governance

Gazebo is widely used for robotics research, allowing physically realistic simulation of sensors, actuators and environments. It is increasingly used to prototype autonomy in safety-critical domains: healthcare, disaster response, transportation, industrial automation.

However, typical Gazebo experiments focus on **task performance** (e.g., throughput, latency, success rate) rather than **ethical performance** (e.g., fairness, rights compliance, harm minimization under constraints). Without a governance layer, reinforcement learning or planning agents may discover **unethical but reward-maximizing behaviors**.

2.3 Medical triage as a stress test

Emergency department triage exposes several hard ethical features:

- **Scarcity**: limited ICU beds, ventilators or staff time.
- **High stakes**: decisions affect survival and long-term outcomes.
- **Fairness and rights**: non-discrimination requirements, obligations to the worst-off.
- **Time pressure**: decisions must be made in seconds, not minutes.

A Gazebo-based **Medical Triage Emergency** scenario therefore provides an ideal testbed for DEME: it combines real-time control, complex moral trade-offs, and strong regulatory expectations.

3. System Overview

The DEME–ErisML Governance Plugin architecture comprises four main components:

1. Gazebo Simulation Environment

- Models patients, staff, equipment, and the triage agent (robot or software agent).
- Advances the physics, sensors and world state at a fixed step.

2. DEME–ErisML Governance Plugin (Gazebo plugin)

- Runs inside Gazebo as a world or model plugin.
- Intercepts proposed actions from the triage agent.
- Extracts relevant state and constructs **EthicalFacts**.
- Communicates with the DEME MCP server.
- Enforces vetoes and selections back into the simulation.

3. DEME MCP Server (Moral Control Plane)

- Network service that wraps the ErisML library.
- Hosts one or more **governance profiles** (e.g., hospital_triage_v2).
- Validates profiles, evaluates actions, and produces decision proofs.
- Optionally anchors proofs into a cryptographic ledger.

4. Analytics and Governance Tooling

- Dashboards for scenario outcomes and fairness metrics.
- Profile authoring tools.
- Log analysis for post-hoc audit and incident investigation.

At each decision step, Gazebo acts as the **environment and actuator layer**, ErisML provides the **ethical semantics**, and the DEME MCP server plus plugin form the **enforcement path**.

4. Architecture and Data Flow

4.1 High-level sequence

For each decision cycle in the triage scenario:

1. **Gazebo updates world state** (patient conditions, vitals, bed availability, etc.).
2. The **triage agent** (e.g., an RL policy node) proposes one or more **candidate actions**, such as:
 - “Assign Patient A to ICU bed 1”
 - “Treat Patient C next and defer others”
3. The **Governance Plugin**:

- Extracts relevant state and candidate actions.
- Constructs EthicalFacts for each option (Section 7.2).
- Sends a request to the **DEME MCP server**: EvaluateOptions(profile_id, EthicalFacts[]).

4. The **DEME MCP server**:

- Uses ErisML to map each EthicalFacts instance to a **moral vector**.
- Applies the selected **governance profile**:
 - Filters options that fall into veto regions.
 - Ranks permissible options.
- Builds a **decision proof** with moral vectors, veto flags, chosen option and explanation.
- Returns {allowed_options, forbidden_options, winner, reason_codes, proof_id}.

5. The **Governance Plugin**:

- Enforces vetoes: blocked actions are never applied in the Gazebo world.
- Applies the chosen option (winner) when multiple remain.
- Optionally feeds the outcome back into the triage agent as a constrained reward signal.

6. The **DEME MCP server**:

- Stores the decision proof locally.
- Periodically batches proofs into Merkle trees and anchors roots into a ledger (optional but recommended for study of audit workflows).

4.2 Message schema (conceptual)

A typical EvaluateOptions call (e.g., JSON over gRPC/REST):

```
{
  "profile_id": "hospital_triage_v2",
  "environment_id": "Gazebo_ED_SIM_01",
```

```
"options": [
  {
    "option_id": "assign_A_ICU1",
    "ethical_facts": { ... }
  },
  {
    "option_id": "assign_B_ICU1",
    "ethical_facts": { ... }
  },
  {
    "option_id": "assign_C_ICU1",
    "ethical_facts": { ... }
  }
]
```

The response:

```
{
  "winner": "assign_A_ICU1",
  "forbidden": ["assign_C_ICU1"],
  "moral_vectors": { "assign_A_ICU1": [ ... ], "...": [ ... ] },
  "reason_codes": {
    "assign_C_ICU1": ["EMTALA_VIOLATION", "FAIRNESS_BELOW_BASELINE"]
  },
  "decision_proof_id": "proof_2025-02-15T02:34:23.127Z"
}
```

The plugin treats this response as **authoritative**: any forbidden option is discarded, and the winner (if provided) is executed.

5. DEME MCP Server

5.1 Responsibilities

The DEME MCP server serves as the **central ethical control plane**:

- **Profile Management**
 - Store, version and activate governance profiles (hospital_triage_v2, urban_av_v1, etc.).
 - Validate profiles using DEME's static validator (veto consistency, DAG acyclicity).
- **Evaluation Services**
 - EvaluateOption – evaluate a single option (for simple controllers).
 - EvaluateOptions – evaluate and resolve among multiple options.
 - SimulateProfileChange – re-run decisions under alternative profiles for sensitivity analysis.
- **Logging and Proof Generation**
 - Store full decision proofs.
 - Aggregate proofs into Merkle trees.
 - Anchor Merkle roots into a ledger (filesystem, database, or external blockchain).
- **Monitoring and Analytics**
 - Expose metrics (rates of vetoes, fairness distribution, rights violations avoided).
 - Provide tracebacks for specific decisions on request.

5.2 ErisML integration

ErisML runs as a library inside the MCP server. For each option:

1. Parses EthicalFacts into internal structures.

2. Calls the appropriate **domain Ethics Modules** (e.g., TriageEM) to compute moral vectors.
3. Passes these vectors to the **governance engine** for veto and scalarization.
4. Returns structured outputs plus explanation tags.

Configuration allows multiple **domain modules** and **policy modules** to be active in the same profile.

6. DEME–ErisML Gazebo Plugin

6.1 Plugin type and lifecycle

The plugin is typically implemented as a **Gazebo world plugin** (for multi-agent triage) or a **model plugin** attached to the triage robot/agent. Its lifecycle:

1. Initialization

- Reads configuration from SDF/INI/YAML:
 - mcp_endpoint
 - profile_id
 - mapping from Gazebo models/topics to patient identities and attributes.
- Establishes a persistent connection to the DEME MCP server.

2. Per-update callback

- At each control tick (e.g., 10–50 Hz), the plugin:
 - Reads the agent’s proposed action(s) from ROS topics or Gazebo messages.
 - Constructs EthicalFacts for each option.
 - Sends an evaluation request and waits for a bounded response time.
 - Applies the response: vetoes or selected action.

3. Shutdown

- Closes connections, flushes local logs,

- Optionally writes a summary report of decisions taken during the episode.

6.2 Configuration example

In an SDF world file:

```
<plugin name="deme_erisml_governance" filename="libdeme_erisml_gazebo.so">

<mcp_endpoint>tcp://localhost:8088</mcp_endpoint>

<profile_id>hospital_triage_v2</profile_id>

<agent_topic>/triage_agent/command</agent_topic>

<governed_topic>/triage_agent/governed_command</governed_topic>

<max_eval_latency_ms>20</max_eval_latency_ms>

</plugin>
```

The triage agent publishes desired commands to `/triage_agent/command`; the plugin publishes **ethically governed commands** to `/triage_agent/governed_command` after consulting ErisML.

7. Medical Triage Emergency Scenario

7.1 Environment

The Gazebo world emulates an **emergency department (ED)** with:

- An ICU area with a small number of **critical care beds**.
- A triage desk where a **triage agent** chooses which patient to admit next.
- Multiple patient models (A, B, C, ...) with:
 - Vital signs (HR, BP, O₂), lab results, SOFA scores.
 - Demographic and social markers (age, housing status).
 - Arrival times and event history.

The triage agent's job is to **assign scarce ICU capacity** and schedule interventions.

7.2 EthicalFacts extraction

The plugin (or a helper node) transforms simulation state into EthicalFacts. For each candidate “assign patient X to ICU” action, we compute:

- `physical_harm_if_delayed` – estimated mortality risk if patient waits.
- `expected_benefit` – probability of significant improvement if treated now.
- `disadvantaged_status` – boolean flag for structurally disadvantaged groups (e.g., unhoused, underinsured).
- `vip_flag` – boolean; should **never** justify preferential treatment.
- `arrival_time` – for fairness and queueing rules.
- `epistemic_confidence` – reliability of risk estimates (e.g., presence of complete diagnostics).
- Legal and policy context flags (e.g., EMTALA applicability).

Example (pseudo-JSON):

```
{
  "patient_id": "A",
  "physical_harm_if_delayed": 0.7,
  "expected_benefit": 0.85,
  "disadvantaged_status": true,
  "vip_flag": false,
  "arrival_time": 1700000000.12,
  "epistemic_confidence": 0.72
}
```

7.3 Governance profile for triage

A sample hospital_triage_v2 profile might specify:

- **Veto rules**
 - Forbid any option where a VIP flag influences priority if clinical need is similar.
 - Forbid actions that cause extreme harm when alternatives with substantially lower harm exist.
 - Forbid discrimination across protected attributes.

- **Lexicographic priorities**
1. Minimize expected harm (beneficence, non-maleficence).
 2. Among similarly effective options, maximize fairness to the worst-off.
 3. Maintain procedural legitimacy (follow institutional policies).

The DEME MCP server uses this profile to interpret the moral vectors produced by TriageEM.

7.4 Example decision step

In a particular simulation tick:

- Patient A and C have similar clinical status; A is unhoused; C is a VIP donor.
- One ICU bed is available.
- The agent proposes three options:
 - assign_A_ICU1, assign_B_ICU1, assign_C_ICU1.

ErisML computes moral vectors:

Option	Harm ↓	Rights	Fairness	Legitimacy
assign_A_ICU1	0.20	1.0	0.90	0.92
assign_B_ICU1	0.30	1.0	0.60	0.95
assign_C_ICU1	0.22	0.25	0.15	0.40

The profile:

- **Vetoes** assign_C_ICU1 because rights and fairness fall below thresholds.
- Compares A vs B lexicographically: A wins on harm and fairness.
- Returns winner = assign_A_ICU1 and a reason code like "EMTALA_VIOLATION" for C.

The Gazebo plugin enforces this:

- Override any request to assign C.
- Apply an action to move A's model to the ICU bed.
- Log the governed decision so researchers can inspect trajectories and fairness statistics.

7.5 Evaluating learning agents

If a reinforcement learning agent controls triage decisions:

- During training, the plugin can **mask vetoed actions** so the agent never receives positive reward for them.
 - During deployment, the agent's proposed action is always filtered through the DEME governance layer.
 - Researchers can examine whether different reward functions still converge to ethically acceptable policies under the same profile.
-

8. Safety and Testing Strategy

To use the DEME–ErisML plugin responsibly, we recommend:

1. Unit and property tests

- Verify that the plugin correctly maps Gazebo state → EthicalFacts.
- Ensure vetoed options are never executed, even under communication delays.

2. Scenario regression suites

- Curate canonical triage scenarios:
 - VIP vs disadvantaged patient with similar needs.
 - Many low-risk vs one high-risk patient.
 - Cases with uncertain diagnostics.
- Run them under multiple governance profiles to confirm behavior.

3. Stress testing

- Vary simulation timestep, network latency to MCP, and number of concurrent triage calls.
- Confirm that the plugin fails **safe** (e.g., conservative defaults) if the MCP server is unreachable or times out.

4. Fairness and rights analysis

- Analyze decision logs by patient group, arrival time, and clinical condition.

- Ensure that discrimination is not reintroduced inadvertently by perception or scoring models upstream of DEME.
-

9. Integration with Learning Systems

The governance plugin is designed to be **agent-agnostic**:

- Classical planners, heuristic policies, and RL agents all produce candidate actions in the same way;
- DEME acts as a **safety and governance envelope** around these systems.

Typical patterns:

- **Constrained exploration** – During training, the plugin masks actions that violate the profile, effectively training in a safe action space.
 - **Reward shaping** – The plugin can inject additional negative rewards for near-veto actions to encourage policies that stay comfortably within acceptable regions.
 - **Curriculum learning** – Profiles can be tightened over time (e.g., from utilitarian to strongly fairness-weighted) while observing how policies adapt.
-

10. Deployment Roadmap

Although the current focus is simulation in Gazebo, the architecture is deliberately close to real robotic stacks:

1. Gazebo + ROS2 simulation →
2. Hardware-in-the-loop testing (real sensors, simulated environment) →
3. Physical pilot deployments in non-clinical environments (e.g., triage training centers) →
4. Controlled clinical pilots once regulatory and ethical approvals are obtained.

Because DEME MCP and ErisML are **network services**, the same governance engine can be reused across simulation and deployment: only the **EthicalFacts extraction** and **actuation interface** change.

11. Conclusion

The **DEME–ErisML Governance Plugin for Gazebo** provides a concrete path from abstract ethical theory to governed behavior in robotics simulations. By connecting Gazebo to a DEME MCP server running ErisML, we can:

- Express hospital triage policies as explicit, versioned governance profiles,
- Enforce hard vetoes and principled rankings in real time,
- Study how learning agents behave under ethical constraints,
- Generate rich, auditable logs for ethics research, regulation and safety engineering.

The **Medical Triage Emergency** scenario demonstrates how such a system can operationalize fairness, rights and harm minimization in a domain where these values are non-negotiable. Future work includes publishing reference implementations of the plugin, expanding the library of triage scenarios, and integrating with broader open-source safety and ethics tooling for robotics.